

## Introduction

if you can build a product that makes 100% of the time, do NOT use ML

ML = learn from experience

## Kinds of ML

### ① supervised learning

predicts targets given inputs

"develop model  $f_{\theta}$  that maps input  $x_i$  to a prediction  $f_{\theta}(x_i)$ "

#### a) regression

"How many", "How much" problems

#### b) classification

• look at feature set and predict which class an example belongs to.

easier to do this in terms of probability

• Cross entropy risk/loss function... potential risk of going wrong - probability of being wrong

#### c) tagging

- multi-label classification  $\leftarrow$  classifying classes that are NOT mutually exclusive

#### d) search & ranking

- which result is MOST useful from irrelevant results

• relevance filter

diff

#### e) recommender systems

• emphasis on personalization to specific users in context of recommender systems

• censored feedback  $\rightarrow$  only people who feel strongly about a subject would provide feedback

• feedback loop  $\rightarrow$  product A is bought more b/c recommended

more... ↓

system thinks product A is better life brought more

## 1) Sequence Learning

most methods throw away past data after every discrete output...

for some cases, we want to build a model based on continuous input

eg: sequence of input  $\rightarrow$  sequence of output  $\Leftrightarrow$  seq<sup>2</sup> seq problem

↳ eg: Tagging and parsing audio speech

automatic speech recognition

text to speech

machine translation - maintain alignment of subject/object  
when translating between languages

eg: dialogue problems

## 2) Unsupervised Learning

any kind of vague, open-ended data science problem

eg: clustering

subspace estimation (if linear - principal component analysis)

representation learning

causality and probabilistic graphical models


generative adversarial networks (GAN)

we want agents, not predictive models  $\Leftrightarrow$  considering interactions w/ a model

## Reinforcement Learning

general problem statement where an agent interacts with an environment over a series of timesteps.

@ timestep  $t$ ,



extremely general

① credit assignment problem  $\rightarrow$  determine which actions to blame/reward for an outcome

② partial observability  $\rightarrow$  current observation may not tell you everything about current state

③ at any timestep, deciding whether to exploit best currently known strategy or explore the action space giving up some short term reward in exchange for knowledge.

- when an environment is completely unknown  $\Rightarrow$  RL problem becomes

a Markov Decision Process

- when the state does not depend on the previous actions, we call the problem

a contextual bandit problem

- when no state, just set of available actions w/ initially unknown rewards,

a multi-armed bandit problem

Hebbian Learning Rule  $\rightarrow$  neurons learn through positive reinforcement

$\hookrightarrow$  basis for Rosenblatt's perceptron learning algorithm